

XV6 SYSTEM CALLS 作业报告

软工2班 3017218092 侯雨茜

PART 1: SYSTEM CALL TRACING

第一部分，需要跟踪系统调用，并把他们打印在终端上。

首先，观察 `syscall.h` 文件：

```
1 // System call numbers
2 #define SYS_fork    1
3 #define SYS_exit    2
4 #define SYS_wait    3
5 #define SYS_pipe    4
6 #define SYS_read    5
7 #define SYS_kill    6
8 #define SYS_exec    7
9 #define SYS_fstat   8
10 #define SYS_chdir   9
11 #define SYS_dup    10
12 #define SYS_getpid 11
13 #define SYS_sbrk   12
14 #define SYS_sleep  13
15 #define SYS_uptime 14
16 #define SYS_open   15
17 #define SYS_write  16
18 #define SYS_mknod  17
19 #define SYS_unlink 18
20 #define SYS_link   19
21 #define SYS_mkdir  20
22 #define SYS_close  21
```

文件中定义了系统调用的名称和对应的序号，这些需要在终端上被显示出来。

接着，修改 `syscall.c` 文件。在文件中添加编号和系统调用名称相互对应的数组，然后在 `syscall()` 函数中添加对应的输出命令：

```

1  #include "types.h"
2  #include "defs.h"
3  #include "param.h"
4  #include "memlayout.h"
5  #include "mmu.h"
6  #include "proc.h"
7  #include "x86.h"
8  #include "syscall.h"
9
10 // User code makes a system call with INT T_SYSCALL.
11 // System call number in %eax.
12 // Arguments on the stack, from the user call to the C
13 // library system call function. The saved user %esp points
14 // to a saved program counter, and then the first argument.
15
16 // Fetch the int at addr from the current process.
17 int
18 fetchint(uint addr, int *ip)
19 {
20     struct proc *curproc = myproc();
21
22     if(addr >= curproc->sz || addr+4 > curproc->sz)
23         return -1;
24     *ip = *(int*)(addr);
25     return 0;
26 }
27
28 // Fetch the nul-terminated string at addr from the current process.
29 // Doesn't actually copy the string - just sets *pp to point at it.
30 // Returns length of string, not including nul.
31 int
32 fetchstr(uint addr, char **pp)
33 {
34     char *s, *ep;
35     struct proc *curproc = myproc();
36
37     if(addr >= curproc->sz)
38         return -1;
39     *pp = (char*)addr;
40     ep = (char*)curproc->sz;
41     for(s = *pp; s < ep; s++){
42         if(*s == 0)
43             return s - *pp;
44     }
45     return -1;
46 }
47
48 // Fetch the nth 32-bit system call argument.
49 int

```

```

50 argint(int n, int *ip)
51 {
52     return fetchint((myproc()->tf->esp) + 4 + 4*n, ip);
53 }
54
55 // Fetch the nth word-sized system call argument as a pointer
56 // to a block of memory of size bytes. Check that the pointer
57 // lies within the process address space.
58 int
59 argptr(int n, char **pp, int size)
60 {
61     int i;
62     struct proc *curproc = myproc();
63
64     if(argint(n, &i) < 0)
65         return -1;
66     if(size < 0 || (uint)i >= curproc->sz || (uint)i+size > curproc-
>sz)
67         return -1;
68     *pp = (char*)i;
69     return 0;
70 }
71
72 // Fetch the nth word-sized system call argument as a string pointer.
73 // Check that the pointer is valid and the string is nul-terminated.
74 // (There is no shared writable memory, so the string can't change
75 // between this check and being used by the kernel.)
76 int
77 argstr(int n, char **pp)
78 {
79     int addr;
80     if(argint(n, &addr) < 0)
81         return -1;
82     return fetchstr(addr, pp);
83 }
84
85 extern int sys_chdir(void);
86 extern int sys_close(void);
87 extern int sys_dup(void);
88 extern int sys_exec(void);
89 extern int sys_exit(void);
90 extern int sys_fork(void);
91 extern int sys_fstat(void);
92 extern int sys_getpid(void);
93 extern int sys_kill(void);
94 extern int sys_link(void);
95 extern int sys_mkdir(void);
96 extern int sys_mknod(void);
97 extern int sys_open(void);

```

```
98 extern int sys_pipe(void);
99 extern int sys_read(void);
100 extern int sys_sbrk(void);
101 extern int sys_sleep(void);
102 extern int sys_unlink(void);
103 extern int sys_wait(void);
104 extern int sys_write(void);
105 extern int sys_uptime(void);
106
107 static int (*syscalls[])(void) = {
108     [SYS_fork]    sys_fork,
109     [SYS_exit]    sys_exit,
110     [SYS_wait]    sys_wait,
111     [SYS_pipe]    sys_pipe,
112     [SYS_read]    sys_read,
113     [SYS_kill]    sys_kill,
114     [SYS_exec]    sys_exec,
115     [SYS_fstat]   sys_fstat,
116     [SYS_chdir]   sys_chdir,
117     [SYS_dup]     sys_dup,
118     [SYS_getpid]  sys_getpid,
119     [SYS_sbrk]    sys_sbrk,
120     [SYS_sleep]   sys_sleep,
121     [SYS_uptime]  sys_uptime,
122     [SYS_open]    sys_open,
123     [SYS_write]   sys_write,
124     [SYS_mknod]   sys_mknod,
125     [SYS_unlink]  sys_unlink,
126     [SYS_link]    sys_link,
127     [SYS_mkdir]   sys_mkdir,
128     [SYS_close]   sys_close,
129 };
130
131 // 用数组存放系统调用对应名字
132 static char syscalls_name[][6] = {
133     [SYS_fork]    "fork",
134     [SYS_exit]    "exit",
135     [SYS_wait]    "wait",
136     [SYS_pipe]    "pipe",
137     [SYS_read]    "read",
138     [SYS_kill]    "kill",
139     [SYS_exec]    "exec",
140     [SYS_fstat]   "fstat",
141     [SYS_chdir]   "chdir",
142     [SYS_dup]     "dup",
143     [SYS_getpid]  "getpid",
144     [SYS_sbrk]    "sbrk",
145     [SYS_sleep]   "sleep",
146     [SYS_uptime]  "uptime",
```

```

147     [SYS_open]      "open",
148     [SYS_write]     "write",
149     [SYS_mknod]     "mknod",
150     [SYS_unlink]    "unlink",
151     [SYS_link]      "link",
152     [SYS_mkdir]     "mkdir",
153     [SYS_close]     "close",
154 };
155
156 void
157 syscall(void)
158 {
159     int num;
160     struct proc *curproc = myproc();
161
162     // 获取系统调用号
163     num = curproc->tf->eax;
164     if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
165         //curproc->tf->eax = syscalls[num]();
166         // 存储系统调用返回值
167         curproc->tf->eax = syscalls[num]();
168         // 打印系统调用名称和编号并换行，使用制表符对齐格式
169         cprintf("\tsyscalls: %s\tid: %d\n", syscalls_name[num], num);
170     } else {
171         cprintf("%d %s: unknown sys call %d\n",
172             curproc->pid, curproc->name, num);
173         curproc->tf->eax = -1;
174     }
175 }

```

最后，在ubuntu中编译运行：

```
1 make qemu-nox
```

运行结果如下：

```

记录了336+1 的读入
记录了336+1 的写出
172124 bytes (172 kB, 168 KiB) copied, 0.000625019 s, 275 MB/s
qemu-system-i386 -nographic -drive file=fs.img,index=1,media=disk,format=raw
-drive file=xv6.img,index=0,media=disk,format=raw -smp 1 -m 512 -snapshot
xv6...
cpu0: starting 0
sb: size 20000 nblocks 19937 ninodes 200 nlog 30 logstart 2 inodestart 32 bm
ap start 58
main-loop: WARNING: I/O thread spun for 1000 iterations
      syscalls: exec id: 7
      syscalls: open id: 15
      syscalls: mknod id: 17
      syscalls: open id: 15
      syscalls: dup id: 10
      syscalls: dup id: 10
i      syscalls: write id: 16
n      syscalls: write id: 16
i      syscalls: write id: 16
t      syscalls: write id: 16
:      syscalls: write id: 16
      syscalls: write id: 16
s      syscalls: write id: 16
t      syscalls: write id: 16
a      syscalls: write id: 16
r      syscalls: write id: 16
t      syscalls: write id: 16
i      syscalls: write id: 16
n      syscalls: write id: 16
g      syscalls: write id: 16
      syscalls: write id: 16
s      syscalls: write id: 16
h      syscalls: write id: 16

      syscalls: write id: 16
      syscalls: fork id: 1
      syscalls: exec id: 7
      syscalls: open id: 15
      syscalls: close id: 21
$      syscalls: write id: 16
      syscalls: write id: 16

```

PART 2: DATE SYSTEM CALL

第二部分，需要添加一个系统调用函数，使其可以返回UTC时间。

`lapic.c` 中定义了函数 `cmostime()` 来读取当前时间。

```

1 // qemu seems to use 24-hour GWT and the values are BCD encoded
2 void
3 cmostime(struct rtcdate *r)
4 {

```

```

5   struct rtcdate t1, t2;
6   int sb, bcd;
7
8   sb = cmos_read(CMOS_STATB);
9
10  bcd = (sb & (1 << 2)) == 0;
11
12  // make sure CMOS doesn't modify time while we read it
13  for(;;) {
14      fill_rtcdate(&t1);
15      if(cmos_read(CMOS_STATA) & CMOS_UIP)
16          continue;
17      fill_rtcdate(&t2);
18      if(memcmp(&t1, &t2, sizeof(t1)) == 0)
19          break;
20  }
21
22  // convert
23  if(bcd) {
24      #define CONV(x)      (t1.x = ((t1.x >> 4) * 10) + (t1.x & 0xf))
25      CONV(second);
26      CONV(minute);
27      CONV(hour  );
28      CONV(day   );
29      CONV(month );
30      CONV(year  );
31      #undef CONV
32  }
33
34  *r = t1;
35  r->year += 2000;
36  }

```

`date.h` 中定义了 `struct rtcdate` 结构体，作为 `cmostime()` 的参数。

```

1  struct rtcdate {
2      uint second;
3      uint minute;
4      uint hour;
5      uint day;
6      uint month;
7      uint year;
8  };

```

添加系统调用函数 `date()`：

1. 在 `syscall.h` 中增加日期的系统调用编号

```

1 // System call numbers
2 #define SYS_fork    1
3 #define SYS_exit    2
4 #define SYS_wait    3
5 #define SYS_pipe    4
6 #define SYS_read    5
7 #define SYS_kill    6
8 #define SYS_exec    7
9 #define SYS_fstat    8
10 #define SYS_chdir   9
11 #define SYS_dup     10
12 #define SYS_getpid  11
13 #define SYS_sbrk    12
14 #define SYS_sleep   13
15 #define SYS_uptime  14
16 #define SYS_open    15
17 #define SYS_write   16
18 #define SYS_mknod   17
19 #define SYS_unlink  18
20 #define SYS_link    19
21 #define SYS_mkdir   20
22 #define SYS_close   21
23 // 日期系统调用编号
24 #define SYS_date    22

```

2. 在 `syscall.c` 中添加系统调用函数的外部声明

```

1 // Fetch the nth word-sized system call argument as a string pointer.
2 // Check that the pointer is valid and the string is nul-terminated.
3 // (There is no shared writable memory, so the string can't change
4 // between this check and being used by the kernel.)
5 int
6 argstr(int n, char **pp)
7 {
8     int addr;
9     if(argint(n, &addr) < 0)
10         return -1;
11     return fetchstr(addr, pp);
12 }
13
14 extern int sys_chdir(void);
15 extern int sys_close(void);
16 extern int sys_dup(void);
17 extern int sys_exec(void);
18 extern int sys_exit(void);
19 extern int sys_fork(void);
20 extern int sys_fstat(void);
21 extern int sys_getpid(void);

```



```
22 extern int sys_kill(void);
23 extern int sys_link(void);
24 extern int sys_mkdir(void);
25 extern int sys_mknod(void);
26 extern int sys_open(void);
27 extern int sys_pipe(void);
28 extern int sys_read(void);
29 extern int sys_sbrk(void);
30 extern int sys_sleep(void);
31 extern int sys_unlink(void);
32 extern int sys_wait(void);
33 extern int sys_write(void);
34 extern int sys_uptime(void);
35 extern int sys_date(void);
36
37 static int (*syscalls[])(void) = {
38     [SYS_fork]    sys_fork,
39     [SYS_exit]    sys_exit,
40     [SYS_wait]    sys_wait,
41     [SYS_pipe]    sys_pipe,
42     [SYS_read]    sys_read,
43     [SYS_kill]    sys_kill,
44     [SYS_exec]    sys_exec,
45     [SYS_fstat]   sys_fstat,
46     [SYS_chdir]   sys_chdir,
47     [SYS_dup]     sys_dup,
48     [SYS_getpid]  sys_getpid,
49     [SYS_sbrk]    sys_sbrk,
50     [SYS_sleep]   sys_sleep,
51     [SYS_uptime]  sys_uptime,
52     [SYS_open]    sys_open,
53     [SYS_write]   sys_write,
54     [SYS_mknod]   sys_mknod,
55     [SYS_unlink]  sys_unlink,
56     [SYS_link]    sys_link,
57     [SYS_mkdir]   sys_mkdir,
58     [SYS_close]   sys_close,
59     [SYS_date]    SYS_date,
60 };
61
62 // 用数组存放系统调用对应名字
63 static char syscalls_name[][6] = {
64     [SYS_fork]    "fork",
65     [SYS_exit]    "exit",
66     [SYS_wait]    "wait",
67     [SYS_pipe]    "pipe",
68     [SYS_read]    "read",
69     [SYS_kill]    "kill",
70     [SYS_exec]    "exec",
```

```
71     [SYS_fstat]    "fstat",
72     [SYS_chdir]    "chdir",
73     [SYS_dup]      "dup",
74     [SYS_getpid]   "getpid",
75     [SYS_sbrk]     "sbrk",
76     [SYS_sleep]    "sleep",
77     [SYS_uptime]   "uptime",
78     [SYS_open]     "open",
79     [SYS_write]    "write",
80     [SYS_mknod]    "mknod",
81     [SYS_unlink]   "unlink",
82     [SYS_link]     "link",
83     [SYS_mkdir]    "mkdir",
84     [SYS_close]    "close",
85     [SYS_date]     "date",
86 };
```

3. 在 `user.h` 中添加用户态函数的定义

```
1  // system calls
2  int fork(void);
3  int exit(void) __attribute__((noreturn));
4  int wait(void);
5  int pipe(int*);
6  int write(int, const void*, int);
7  int read(int, void*, int);
8  int close(int);
9  int kill(int);
10 int exec(char*, char**);
11 int open(const char*, int);
12 int mknod(const char*, short, short);
13 int unlink(const char*);
14 int fstat(int fd, struct stat*);
15 int link(const char*, const char*);
16 int mkdir(const char*);
17 int chdir(const char*);
18 int dup(int);
19 int getpid(void);
20 char* sbrk(int);
21 int sleep(int);
22 int uptime(void);
23 int date(struct rtcdate*); // 日期
```

4. 在 `usys.s` 中添加用户态函数的实现

```
1  SYSCALL(fork)
2  SYSCALL(exit)
```

```
3  SYSCALL(wait)
4  SYSCALL(pipe)
5  SYSCALL(read)
6  SYSCALL(write)
7  SYSCALL(close)
8  SYSCALL(kill)
9  SYSCALL(exec)
10 SYSCALL(open)
11 SYSCALL(mknod)
12 SYSCALL(unlink)
13 SYSCALL(fstat)
14 SYSCALL(link)
15 SYSCALL(mkdir)
16 SYSCALL(chdir)
17 SYSCALL(dup)
18 SYSCALL(getpid)
19 SYSCALL(sbrk)
20 SYSCALL(sleep)
21 SYSCALL(uptime)
22 SYSCALL(date)
```

5. 在 `sysproc.c` 中添加系统调用函数 `sys_date` 的实现

```
1  // return UTC time
2  int
3  sys_date(struct rtcdate *r)
4  {
5      if (argptr(0, (void *)&r, sizeof(*r)) < 0) {
6          return -1;
7      }
8      cmostime(r);    // get time from cmos
9      return 0;
10 }
```

6. 新建 `date.c` 文件，添加使用系统调用函数 `sys_date` 的方法

```
1  #include "types.h"
2  #include "user.h"
3  #include "date.h"
4
5  int
6  main(int argc, char *argv[])
7  {
8      struct rtcdate r;
9
10     if (date(&r)) {
11         printf(2, "date failed\n");
```

```

12     exit();
13 }
14
15 // your code to print the time in any format you like...
16 printf(1, "%d-%d %d %d:%d:%d\n", r.month, r.day, r.year, r.hour,
r.minute, r.second);
17     exit();
18 }

```

7. 在 `Makefile` 中添加 `UPROGS` 对应命令的定义

```

1  UPROGS=\
2  _cat\
3  _echo\
4  _forktest\
5  _grep\
6  _init\
7  _kill\
8  _ln\
9  _ls\
10 _mkdir\
11 _rm\
12 _sh\
13 _stressfs\
14 _usertests\
15 _wc\
16 _zombie\
17 _big\
18 _date\
19
20 fs.img: mkfs README $(UPROGS)
21     ./mkfs fs.img README $(UPROGS)

```

编译并运行，输入命令 `date`，日期无法以正确格式显示。推测是由于Part 1中的输出影响了时间的显示。为防止Part 1的输出影响时间的显示，将 `syscall.c` 中的一些代码注释掉：

```

1  #include "types.h"
2  #include "defs.h"
3  #include "param.h"
4  #include "memlayout.h"
5  #include "mmu.h"
6  #include "proc.h"
7  #include "x86.h"
8  #include "syscall.h"
9
10 // User code makes a system call with INT T_SYSCALL.
11 // System call number in %eax.

```

```

12 // Arguments on the stack, from the user call to the C
13 // library system call function. The saved user %esp points
14 // to a saved program counter, and then the first argument.
15
16 // Fetch the int at addr from the current process.
17 int
18 fetchint(uint addr, int *ip)
19 {
20     struct proc *curproc = myproc();
21
22     if(addr >= curproc->sz || addr+4 > curproc->sz)
23         return -1;
24     *ip = *(int*)(addr);
25     return 0;
26 }
27
28 // Fetch the nul-terminated string at addr from the current process.
29 // Doesn't actually copy the string - just sets *pp to point at it.
30 // Returns length of string, not including nul.
31 int
32 fetchstr(uint addr, char **pp)
33 {
34     char *s, *ep;
35     struct proc *curproc = myproc();
36
37     if(addr >= curproc->sz)
38         return -1;
39     *pp = (char*)addr;
40     ep = (char*)curproc->sz;
41     for(s = *pp; s < ep; s++){
42         if(*s == 0)
43             return s - *pp;
44     }
45     return -1;
46 }
47
48 // Fetch the nth 32-bit system call argument.
49 int
50 argint(int n, int *ip)
51 {
52     return fetchint((myproc()->tf->esp) + 4 + 4*n, ip);
53 }
54
55 // Fetch the nth word-sized system call argument as a pointer
56 // to a block of memory of size bytes. Check that the pointer
57 // lies within the process address space.
58 int
59 argptr(int n, char **pp, int size)
60 {

```

```

61     int i;
62     struct proc *curproc = myproc();
63
64     if(argint(n, &i) < 0)
65         return -1;
66     if(size < 0 || (uint)i >= curproc->sz || (uint)i+size > curproc-
>sz)
67         return -1;
68     *pp = (char*)i;
69     return 0;
70 }
71
72 // Fetch the nth word-sized system call argument as a string pointer.
73 // Check that the pointer is valid and the string is nul-terminated.
74 // (There is no shared writable memory, so the string can't change
75 // between this check and being used by the kernel.)
76 int
77 argstr(int n, char **pp)
78 {
79     int addr;
80     if(argint(n, &addr) < 0)
81         return -1;
82     return fetchstr(addr, pp);
83 }
84
85 extern int sys_chdir(void);
86 extern int sys_close(void);
87 extern int sys_dup(void);
88 extern int sys_exec(void);
89 extern int sys_exit(void);
90 extern int sys_fork(void);
91 extern int sys_fstat(void);
92 extern int sys_getpid(void);
93 extern int sys_kill(void);
94 extern int sys_link(void);
95 extern int sys_mkdir(void);
96 extern int sys_mknod(void);
97 extern int sys_open(void);
98 extern int sys_pipe(void);
99 extern int sys_read(void);
100 extern int sys_sbrk(void);
101 extern int sys_sleep(void);
102 extern int sys_unlink(void);
103 extern int sys_wait(void);
104 extern int sys_write(void);
105 extern int sys_uptime(void);
106 extern int sys_date(void);
107
108 static int (*syscalls[])(void) = {

```

```
109 [SYS_fork]    sys_fork,
110 [SYS_exit]    sys_exit,
111 [SYS_wait]    sys_wait,
112 [SYS_pipe]    sys_pipe,
113 [SYS_read]    sys_read,
114 [SYS_kill]    sys_kill,
115 [SYS_exec]    sys_exec,
116 [SYS_fstat]   sys_fstat,
117 [SYS_chdir]   sys_chdir,
118 [SYS_dup]     sys_dup,
119 [SYS_getpid]  sys_getpid,
120 [SYS_sbrk]    sys_sbrk,
121 [SYS_sleep]   sys_sleep,
122 [SYS_uptime]  sys_uptime,
123 [SYS_open]    sys_open,
124 [SYS_write]   sys_write,
125 [SYS_mknod]   sys_mknod,
126 [SYS_unlink]  sys_unlink,
127 [SYS_link]    sys_link,
128 [SYS_mkdir]   sys_mkdir,
129 [SYS_close]   sys_close,
130 [SYS_date]    sys_date,
131 };
132
133 // 用数组存放系统调用对应名字
134 //static char syscalls_name[][6] = {
135 //    [SYS_fork]    "fork",
136 //    [SYS_exit]    "exit",
137 //    [SYS_wait]    "wait",
138 //    [SYS_pipe]    "pipe",
139 //    [SYS_read]    "read",
140 //    [SYS_kill]    "kill",
141 //    [SYS_exec]    "exec",
142 //    [SYS_fstat]   "fstat",
143 //    [SYS_chdir]   "chdir",
144 //    [SYS_dup]     "dup",
145 //    [SYS_getpid]  "getpid",
146 //    [SYS_sbrk]    "sbrk",
147 //    [SYS_sleep]   "sleep",
148 //    [SYS_uptime]  "uptime",
149 //    [SYS_open]    "open",
150 //    [SYS_write]   "write",
151 //    [SYS_mknod]   "mknod",
152 //    [SYS_unlink]  "unlink",
153 //    [SYS_link]    "link",
154 //    [SYS_mkdir]   "mkdir",
155 //    [SYS_close]   "close",
156 //    [SYS_date]    "date",
157 //};
```

```

158
159 void
160 syscall(void)
161 {
162     int num;
163     struct proc *curproc = myproc();
164
165     // 获取系统调用号
166     num = curproc->tf->eax;
167     if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
168         //curproc->tf->eax = syscalls[num]();
169         // 存储系统调用返回值
170         curproc->tf->eax = syscalls[num]();
171         // 打印系统调用名称和编号并换行, 使用制表符对齐格式
172         //cprintf("\tsyscalls: %s\tid: %d\n", syscalls_name[num], num);
173     } else {
174         cprintf("%d %s: unknown sys call %d\n",
175             curproc->pid, curproc->name, num);
176         curproc->tf->eax = -1;
177     }
178 }

```

编译并运行, 输入命令 `date`, 日期以正确的格式显示:

```

cpu0: starting 0
sb: size 20000 nblocks 19937 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap s
tart 58
init: starting sh
$ date
10-22 2019 2:0:49
$

```